



MicroPhonon Acoustic Leak Detector v0.1

Principle of Operation

When pressurized water escapes from a seam, crack, or loose fitting on plumbing, broadband high frequency acoustics are emitted. These signals can travel an extended distance (10m or more) in free-space and be detected by a remote sensor. This suggests a stand-off method for water leak detection that is substantially different compared to conventional contact moisture sensors and various types of smart flow meters installed on or inside of plumbing lines.

The acoustic leak detector combines the principles of glass breakage and smoke detectors. Glass breakage occurs on a timescale of a few seconds or less, requiring continuous data acquisition at rates approaching 1 kHz or more. Almost all plumbing leaks are persistent, which allows for sub-Hz sampling rates and exceptional battery life.

Glass breakage has the advantage of being generally louder than the ambient environment so that detector sensitivity is not an issue. Small water leaks, especially when separated from the sensor by acoustic barriers such as walls, are likely to be at or below the environmental background. Much higher sensitivity and noise mitigation is essential. This can be addressed by active and passive filtering and judicious use of statistical analysis.

The acoustic approach to water leak detection presented here works well but is not capable of identifying all water issues. It will not respond to leaking roofs or melting ice, for example. The nature of the pressurized orifice is also important. A smooth hole drilled into a pipe will produce a Laminar flow jet that emits a much weaker acoustic signal – perhaps undetectable – compared to a jagged crack or loose fitting at the same backing pressure. Acoustic sensing can be used to complement standard leak detection techniques and will be especially useful when monitoring large areas.

The sensor is fabricated on a 48 x 33 mm² PCB using standard, low cost electronic components. It is designed to work with a host controller as part of a low power wireless system with a 1.8–3.6V power source. Two-way communication takes places via I2C. All signal processing and analysis occurs on the sensor; no audio is sent to the cloud. Sensor information is available on-demand as 12 status bytes.

Analog front-end

The acoustic transducer is a surface-mount MEMS microphone that forms one end of a rectangular Helmholtz resonator. The resonator center frequency is ~ 8 kHz with $Q \sim 30$. The microphone signal is amplified 40 dB and sent to a 2-stage, 4th-order Sallen-Key high-pass filter with 3 dB cutoff at 8 kHz. This is followed by a 23 dB post-filter gain stage. The microphone is pulse-biased and the amplifier

stages are gated on-off to increase battery life. To illustrate, if a single audio acquisition occurs every 2 seconds (default), the sensor will be in sleep mode > 99% of the time.

Signal processing

After the filtered signal leaves the analog front-end, a 12-bit ADC in the MCU collects 256 sample points at 33.3 kS/s corresponding to a 7.68 ms duration temporal acquisition. The data is checked for ADC saturation and tagged as noisy if necessary. If no saturation has occurred, the MCU performs an FFT of the time data using a uniform window, as appropriate for broadband white noise. The LEA of the Texas Instruments MSP430FR5994 calculates the FFT > 10x faster and more efficiently than an ARM-Cortex M0+. The spectral data is rendered as an array of complex numbers separated by 130.2 Hz frequency steps. A 34 point subset in the range 7–11.5 kHz is selected. The spectral energy at each frequency step is calculated and then summed to provide the total acoustic energy in this frequency window. The combination of mechanical, analog, and digital filtering greatly reduces the influence of background noise, dramatically reduces the potential of ADC saturation, and eliminates the possibility of intelligible conversations being discerned if the sensor network becomes compromised, i.e. eavesdropping is impossible.

Analysis

Each polling cycle analyzes the acoustic signal and designates it as one of the following three events: quiet, leak, or noise. These events are accumulated in three separate data arrays. The value 1 is entered if that specific event occurs; otherwise it is counted as 0. For example: if a leak signal is detected, a 1 is added to the leak array while 0 is placed in the quiet and noise arrays. As new events are added, the oldest events are removed, i.e. the fixed-size data arrays slide in time.

All data arrays are the same size and contain elements that are either 0 or 1. The array size is configured by the host in the range 10–255; default is 20. Size and polling rate determine the time over which a full data set is acquired. The default 0.5 Hz polling rate and 20 event set size define a 40 second data accumulation window. This accumulation window will depend on the application. If the environment is known to always be nominally quiet, a short time window will identify a potential leak quickly. In a bathroom where a shower may be running 20 minutes or more, the accumulation window must be longer than this to avoid false alarms.

If leak events become an appreciable portion of the accumulation window, an alarm condition may exist. Ideally, all events in the window will be identified as leaks. In practice, environmental noise, irregularities in water flow, pressure fluctuations, the presence of bubbles, or other acoustic disturbances may be recorded as quiet or noise events – not leaks. To account for non-ideal signals, the alarm trigger size can be independently set below the event set size. The trigger value is not critical; testing has shown that 80–90 percent of the event set size is a good choice. The default values are 18 for the trigger size and 20 for the event size.

When a potential leak is identified, the sensor provides an immediate alert by toggling the digital ALARM line. The host controller can monitor this line with an interrupt. The corresponding status byte is also set; see I2C communications below.

The count of quiet, leak, and noise events are available at any time by querying the sensor via I2C. If noise events become significant, the noise status byte is set. This is important, as leak detection is

compromised when ongoing noise is present. The count of quiet, leak, and noise events sum to the event set size.

Noise mitigation

To hear the smallest leaks, the electronics must be configured to have high sensitivity to very weak signals. The trade-off is increased susceptibility to ambient environmental noise that can increase the probability of false alarms. The sensor is adept at identifying various types of noise that it ignores. Excessive noise will, however, mask the presence of a leak. By accumulating statistics from each polling cycle, the sensor can inform the host controller that it is operating in a noisy environment.

On power-up or after a reset, the sensor learns about the ambient acoustic environment. It collects a set of acoustic signals of duration ~ 8 ms at a 1 Hz rate. The size of the background data set can be in the range 10--255; default is 30 corresponding to a 30 second total acquisition time. The data is analyzed to get an average background (\bar{x}_B) and standard deviation (σ_B). If $2\sigma_B > \bar{x}_B$, the background is determined to be not quiet. The red LED blinks twice and a new background data set is collected. To obtain additional noise rejection, the acoustic data stream is checked for ADC saturation. Excessively loud signals are discarded and the individual acquisition is repeated. This continues until the background acoustic environment is sufficiently quiet and stable.

The ambient acoustic level and electronic noise floor determine the sensitivity to leaks. For this reason, a well filtered supply voltage should be placed on the 3V3 terminal. Battery power does not guarantee that noise from other electronics will not propagate into the sensor board and cause problems.

Once the background is established, the sensor begins polling at a user defined rate; default is 1 Hz. Slower polling rates increase battery life. The ADC audio samples are first checked for saturation. If an excessively loud signal is present, a **noise event** is recorded, the red LED blinks twice, and the firmware waits for the next polling cycle.

The training session sets the threshold for detecting a leak signal, which is one standard deviation above the noise floor: $\bar{x}_B + \sigma_B$. If there is no ADC saturation, the spectral energy (x_1) derived from an FFT of the temporal data is checked. If $x_1 < \bar{x}_B + \sigma_B$, no leak is detected and a **quiet event** is recorded; the green LED blinks once. If $x_1 > \bar{x}_B + \sigma_B$, additional processing commences: four sequential acoustic measurements are made, separated by 45 ms. If any of these additional acquisitions saturate the ADC, further processing aborts, an environmental **noise event** is recorded, the red LED blinks twice, and the polling loop waits for the next iteration.

If all five filtered audio signals are within the dynamic range of the ADC, their individual spectral energies (x_i ; $i=1-5$) are used to calculate the average energy (\bar{x}) and standard deviation (σ). If $2\sigma > \bar{x}$, excessive fluctuations are present in the data set. This indicates a level of environmental noise that prevents reliable leak detection. This polling cycle is recorded as a **noise event**, the red LED blinks twice, and the program waits for the next timed iteration.

If the acoustic energy fluctuations are sufficiently small, the next step is to check if all five spectral energies (x_i) exceed the leak threshold: $x_i > \bar{x}_B + \sigma_B$. When this occurs, a **leak event** is recorded and the red LED blinks once. Otherwise, this loop iteration is associated with impulse noise, but recorded as a **quiet event**. The green LED blinks twice.

Performing the above sequence of tests and checks ensures that only a persistent, stable signal in the target acoustic frequency range is counted as a **leak event**. This represents a trade-off between sensitivity and reliability to reduce false alarms.

I2C write and read operations

This section describes how a master controller interfaces with the slave sensor on the I2C bus. Recommended clock rate: 100 kHz. External pullup resistors required (4.7k recommended). The default I2C address 0x77.

1. Commands (write only). All commands are sent as two-byte pairs. Send START, followed by the 7-bit slave address and the R/W bit cleared, then the desired function hex code and a single byte data field as shown in Table 1. Finally, assert NACK and STOP. Between 1 and 8 pairs (function byte + data byte) can be sent, provided there are an even number of bytes. The hex code and/or data byte can be sent as hexadecimal or decimal. All write commands clear the counting arrays.

FUNCTION	HEX CODE	DATA BYTE
System reset	0x72	0x72
Alarm	0x61	0x30 or 0x31 (Off or On) Default Off
Sleep	0x70	0x30 or 0x31 (Sleep or Active) Default Active
Background array size	0x62	0x0A – 0xFF (10 – 255) Default: 30
Event array size	0x65	0x0A – 0xFF (10 – 255) Default: 20
Alarm trigger size	0x74	Less than or equal to event array size; Default: 18
Polling period	0x6F	0x01 – 0x09 (1 – 30 sec) Default: 1 sec; see below
LED	0x6C	0x30 or 0x31 (Off or On) Default On

Table 1. Two-byte write commands.

The function descriptions are as follows:

- System reset restores the default settings. This can be sent with other functions, but they will be ignored.
- An alarm state is cleared with 0x61 0x30. Alarm can be set by master with 0x61 0x31; use for testing.
- Sending 0x70 0x30 turns the sensor off and reduces power consumption to minimum. Previously set monitor functions in Table 1 are retained. The sensor is re-activated with 0x70 0x31.
- Background size: The sensor will automatically train itself to learn about the nominally quiet background, i.e. with no noise or leaks present. The background acoustic level is acquired at 1 second intervals for the specified number of counts. Default background size is 30 counts = 30 seconds. A new training session is performed: i) at startup, ii) whenever the background count is set, eg. 0x62 0x1E, or iii) on system reset. The training session sets the sensitivity level.
- The event array size and alarm trigger size are configurable from 10 to a maximum of 255. These arrays record the number of events below and exceeding the sensitivity level found in the training session. The oldest events are discarded from the arrays as each new event is added. An alarm condition exists when the sum of events above the sensitivity level (leak_count) reaches the alarm trigger size. Alarm trigger size must be less than or equal to event array size.

- Polling period determines the rate at which the sensor makes an acoustic measurement. The period is set with one of the data bytes in Table 2. Longer measurement intervals reduce battery power consumption.

0x01	0x02	0x03	0x04	0x05	0x06	0x07	0x08	0x09
1 sec	2 sec	3 sec	5 sec	10 sec	15 sec	20 sec	25 sec	30 sec

Table 2. Polling period data byte is set with function command 0x25

- The LED function (0x6C) disables (0x30) or enables (0x31) surface-mount LEDs on the PCB. Visual indicators can be useful for configuring and evaluating the device, but should be turned off when the sensor is in operation. Default is enabled (on).

Example command string: The 4-byte sequence 0x61 0x30 0x70 0x31 clears an alarm state and resumes monitoring.

Monitor parameters can be set while in the sleep state.

2. Sensor status (read only). The sensor status can be obtained at any time by asserting START, sending the 7-bit slave address with the R/W bit set, and then reading the 12-byte sequence shown in Table 3. The master should follow the sequential read by asserting NACK then STOP.

Byte 0	Byte 1	Byte 2	Byte 3	Byte 4	Byte 5
Alarm	Noise alert	quiet_count	leak_count	noise_count	background_size
Byte 6	Byte 7	Byte 8	Byte 9	Byte 10	Byte 11
event_size	trigger_size	Period	LED	Sleep	Version

Table 3. 12-byte sensor status string

- Byte 0 is set to 0x01 if an alarm exists; 0x00 otherwise. The digital output line marked ALARM is toggled when an alarm situation is detected. The host can monitor this line with an interrupt, so it is not necessary to poll the status of Byte 0. The firmware can be configured to place the sensor into a sleep state when an alarm is detected or it can continue acquiring and processing data.
- Byte 1 is set when $\text{leak_count} < \text{trigger_size}$ **AND** $\text{leak_count} + \text{noise_count} \geq \text{trigger_size}$; 0x00 otherwise. This provides a warning that the sensor is operating in a noisy environment.
- Byte 2 is the current count (in hexadecimal) of below-sensitivity level events.
- Byte 3 is the current count (in hexadecimal) of above-sensitivity level events.
- Byte 4 is the count (in hexadecimal) of noise events. The sensor ignores these events when assessing for the presence of a leak, which helps prevent false alarms.
- Bytes 5–8 are defined in Section 1.
- Byte 9 indicates if the LEDs are off (0x00) or on (0x01).

- Byte 10 indicates if the sensor is sleeping (0x00) or active (0x01).
- Byte 11 is the firmware version running on the slave.